


```
# System Daignostics
import platform
from datetime import datetime

print("Machine ID: ", platform.node())
print("Timestamp: ", datetime.now())
```

Machine ID: 2025aa05327
Timestamp: 2026-08-07 21:45:36.753165

Screenshot with Timestamp



Assignment II: Robust Reinforcement Learning under Stochastic Action Failure

Course: Reinforcement Learning

Group Contribution Declaration

Member Name	Student ID	Percentage Contribution (%)
Anirban Sen	2025AA05327	20%
Sukanya Yadav	2025AA05630	20%

Member Name	Student ID	Percentage Contribution (%)
Anirban Gupta	2025AA05172	20%
Sneha Singh	2025AA05176	20%
Rohit Gupta	2025AA05359	20%

Objective

This assignment investigates the impact of **stochastic actuator failure** on Reinforcement Learning (RL) agents in the **LunarLander - v3** Gymnasium environment. Control commands in real-world robotics often fail due to hardware faults, unreliability, or communication delays. We design a modified environment to model these real-world constraints and evaluate the resilience and performance of **Deep Q-Networks (DQN)** versus **Double Deep Q-Networks (DDQN)**.

Key Components & Tasks

- **Task (a): Modified Environment Implementation (RobustLunarLanderWrapper)**
 - **Intermittent Engine Misfire:** Simulated a stochastic failure rate where attempted thruster actions ($a \in \{1, 2, 3\}$) are randomly replaced with $a=0$ ("Do Nothing") with a probability of $P=0.15$.
 - **Action-Intent Fuel Penalty:** Applied a -0.3 penalty based on the agent's *intended* action rather than the executed action, penalizing attempted fuel consumption even on misfires.
 - **Safe Landing Bonus:** Awarded a $+50.0$ bonus strictly when five simultaneous criteria are met: episode termination without truncation, both landing gear contacts active (`obs[6] == 1`, `obs[7] == 1`), and strict bounds on horizontal velocity ($|v_x| < 0.10$), vertical velocity ($|v_y| < 0.10$), and tilt angle ($|\theta| < 0.10$).
 - **Information Shielding:** Ensured failure metadata is masked in the output `info` dictionary to test the agents' capacity to learn under unobserved stochasticity.
- **Task (b): Deep Q-Network (DQN)**
 - Implemented standard DQN featuring an Experience Replay Buffer (capacity: 100,000, batch size: 64), an ϵ -greedy exploration schedule ($1.0 \rightarrow 0.01$, decay 0.995), and soft target network updates ($\tau=0.005$).
- **Task (c): Double Deep Q-Network (DDQN)**
 - Implemented DDQN to mitigate Q-value overestimation bias by decoupling action selection (using the online network) from evaluation (using the target network).
- **Task (d): Performance Evaluation**
 - Generated 4 comparative plots across 5 random seeds: Episode Rewards, Moving Average Rewards (100 episodes), TD Loss, and Q-Value Trajectories.
- **Task (e): Discussion & Analysis**

- Evaluated overall performance, overestimation bias mitigation, policy adaptiveness under failure, and potential extensions like Prioritized Experience Replay (PER) or Dueling architectures.

1. Setup and Environment Initialization

```
#
=====

# Setup & Dependencies Installation
#
=====

# Install required packages quietly (-q flag suppresses unnecessary
output log)
# Note: 'gymnasium[box2d]' installs Swig and the Box2D engine required
for LunarLander
!pip install gymnasium[box2d] torch numpy matplotlib -q

# Standard Library Imports (type hints, utilities)
import random
import time
from typing import Tuple, List, Dict

# Environment & Numerical Computing Libraries
import gymnasium as gym
import numpy as np
import matplotlib.pyplot as plt

# PyTorch Deep Learning Modules
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

#
=====

# Reproducibility Configuration
#
=====

# Set a global seed across all pseudorandom generators to ensure
# deterministic model behavior and reproducible experimental results
RANDOM_SEED = 42

random.seed(RANDOM_SEED)
np.random.seed(RANDOM_SEED)
```

```

torch.manual_seed(RANDOM_SEED)

# Set seed for CUDA operations if training on a GPU
if torch.cuda.is_available():
    torch.cuda.manual_seed(RANDOM_SEED)

#
=====
=====
# Compute Device Configuration
#
=====
=====

# Automatically detect and assign CUDA GPU acceleration if present;
default to CPU
DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {DEVICE}")

===== 0.0/3.7 MB ? eta -:--:--
===== 0.8/3.7 MB 25.4 MB/s eta
0:00:01 ----- 3.7/3.7 MB 71.8 MB/s
eta 0:00:01 ----- 3.7/3.7 MB 52.6
MB/s eta 0:00:00
===== 1.9/1.9 MB 91.1 MB/s eta
0:00:00
===== 12.8/12.8 MB 98.5 MB/s eta
0:00:00

```

Task (a): Environment Overview & Implementation Rules

Before constructing custom environment wrappers, engine failure models, and modified reward structures, we must establish a clear baseline for how **LunarLander-v3** operates out of the box.

At its core, **LunarLander** is a classic continuous-state, discrete-action control problem from the Gymnasium Box2D physics environment family. The objective is to safely pilot a lunar module down onto a designated landing pad anchored between two flags at the origin ($x=0, y=0$).

1. The Observation Space (8 Continuous State Variables)


At every timestep t , the environment returns an 8-dimensional state vector $s_t \in R^8$:

Index	Symbol	Parameter	Description
obs[0]	x	Horizontal Position	Centered at 0.0 (landing pad center). Range

Index	Symbol	Parameter	Description
obs[1]	y	Vertical Position	roughly [-1, 1] . Surface is near 0.0 ; starts near top of screen (≈ 1.4).
obs[2]	v_x	Linear Horizontal Velocity	Speed moving left/right.
obs[3]	v_y	Linear Vertical Velocity	Speed moving up/down (negative = falling).
obs[4]	θ	Angle / Tilt	Orientation in radians (0.0 = perfectly upright).
obs[5]	ω	Angular Velocity	Rate of rotation.
obs[6]	l	Left Leg Ground Contact	Boolean trigger (1.0 if touching ground, else 0.0).
obs[7]	r	Right Leg Ground Contact	Boolean trigger (1.0 if touching ground, else 0.0).

2. The Action Space (4 Discrete Controls)

The agent controls the lander using 4 discrete control choices ($a \in \{0, 1, 2, 3\}$):



- **0: Do Nothing** — Coast passively under lunar gravity.
- **1: Fire Left Orientation Engine** — Pushes the craft right / tilts counter-clockwise.
- **2: Fire Main Engine** — Provides upward vertical thrust.
- **3: Fire Right Orientation Engine** — Pushes the craft left / tilts clockwise.

3. Standard Environment vs. Task (a) Modification

Standard LunarLander - v3:

- **Deterministic Actions:** Executing action **2** fires the main thruster **100%** of the time.

- **Default Rewards:** Points are awarded for approaching the pad and penalized for moving away. Crashing incurs a **-100** penalty, while a successful touchdown yields **+100**.

Task (a) Modifications:

Task (a) forces the agent to operate under simulated real-world hardware failure and stricter flight constraints:

- **Stochastic Hardware Misfires:** Thruster actions (1, 2, 3) suffer a **15%** random misfire rate, overriding the action to **0** (Do Nothing).
 - **Intent-Based Fuel Penalties:** A penalty of **-0.3** is deducted whenever the agent *attempts* to fire a thruster, even if an engine misfire occurs.
 - **Strict Safe Landing Criteria:** A high-value **+50** bonus is awarded only if the lander meets strict speed, angle, and leg-contact requirements at termination.
-

4. Implementation Rules for RobustLunarLanderWrapper

When implementing the custom environment wrapper, the `step` function must follow this exact order of execution:

- **Step 1:** Store the original action a selected by the agent policy.
- **Step 2:** Check for misfires. If $a \in \{1, 2, 3\}$ and a random uniform float $r < 0.15$, replace a with $a_{\text{exec}} = 0$ (Do Nothing). If $a = 0$, execute $a_{\text{exec}} = 0$ without modification.
- **Step 3:** Pass and execute a_{exec} in the base `LunarLander-v3` environment to obtain `obs`, `R_base`, `terminated`, `truncated`, and `info`.
- **Step 4:** Compute the modified reward function:

$$R = R_{\text{base}} - 0.3 \times 1(a \neq 0) + B$$

(where $1(a \neq 0)$ is an indicator function that equals 1 if the agent intended to use a thruster, and 0 otherwise).

- **Step 5:** Compute the landing bonus $B = 50$ if all 5 conditions hold simultaneously:
 - `terminated == True`
 - `truncated == False`
 - `obs[6] == 1.0` (left leg contact)
 - `obs[7] == 1.0` (right leg contact)
 - Kinematic thresholds met: $|\text{obs}[2]| < 0.10$ (v_x), $|\text{obs}[3]| < 0.10$ (v_y), and $|\text{obs}[4]| < 0.10$ (angle θ).
- **Step 6:** Return the modified step tuple (`obs`, `R`, `terminated`, `truncated`, `info`) ensuring all internal misfire indicators are scrubbed from the `info` dictionary to prevent state-leakage.

```
class RobustLunarLanderWrapper(gym.Wrapper):
    """
    Custom Gymnasium Wrapper for LunarLander-v3.
```

Modifies the base environment to simulate real-world physical constraints and actuator failure:

1. Intermittent Thruster Misfire: Thruster actions (1, 2, 3) randomly fail with probability `failure\_prob`.

*2. Action-Intent Fuel Penalty: Deducts 0.3 points whenever the agent *intends* to use a thruster ($a \neq 0$).*

3. Precise Safe Landing Bonus: Awards +50 points if strict physical criteria are met upon termination.

4. Information Scrubbing: Cleans internal failure indicators from the returned info dictionary.

```
"""

def __init__(
    self,
    env: gym.Env,
    failure_prob: float = 0.15,
    landing_bonus: float = 50.0,
):
    """
    Initialize the environment wrapper.

    Args:
        env (gym.Env): Base LunarLander-v3 environment instance.
        failure_prob (float): Probability of a thruster misfire
        occurring (default: 0.15 or 15%).
        landing_bonus (float): Bonus reward awarded for a precise
        safe landing (default: +50.0).
    """
    super().__init__(env)
    self.failure_prob = failure_prob
    self.landing_bonus = landing_bonus

    # Verification metrics to track empirical misfire rates across
    training steps
    self.total_thruster_attempts = 0
    self.total_misfires = 0

    def step(self, action: int):
        """
        Executes one step in the environment with stochastic failure
        and modified rewards.

        Args:
            action (int): The original action chosen by the agent's
            policy (0, 1, 2, or 3).

        Returns:
            Tuple: (obs, modified_reward, terminated, truncated,
            clean_info)
        """
```

```

# Step 1: Store Original Action Intent
original_action = int(action)

# Step 2: Simulate Intermittent Engine Failure (Stochasticity)
a_exec = original_action

# Actions 1 (Left), 2 (Main), and 3 (Right) correspond to
thruster engines
if original_action in {1, 2, 3}:
    self.total_thruster_attempts += 1

    # Sample uniform random value  $r \sim U(0.0, 1.0)$ 
    r = np.random.uniform(0.0, 1.0)
    if r < self.failure_prob:
        a_exec = 0 # Misfire triggered: Override action to 0
        (Do Nothing)
        self.total_misfires += 1

    elif original_action == 0:
        a_exec = 0 # Action 0 (Do Nothing) executes
        deterministically

    else:
        raise ValueError(
            f"Invalid action {original_action}. Expected action in
            {{0, 1, 2, 3}}."
        )

# Step 3: Execute Transformed Action in Base Environment
obs, base_reward, terminated, truncated, info =
self.env.step(a_exec)

# Step 4: Compute Fuel Penalty Based on Action Intent
# Penalty of -0.3 is charged for ATTEMPTING to fire a thruster
(original_action != 0),
# even if an engine misfire occurs during execution.
fuel_penalty = -0.3 if original_action != 0 else 0.0

# Step 5: Evaluate Safe Landing Bonus Conditions
landing_bonus = 0.0

# Landing bonus is only checked if the episode ended naturally
(not time limit truncation)
if terminated and not truncated:
    # Observation indices mapping:
    # obs[2]: Horizontal velocity ( $v_x$ )
    # obs[3]: Vertical velocity ( $v_y$ )
    # obs[4]: Lander angle/tilt ( $\theta$ )
    # obs[6]: Left leg contact indicator (1.0 = ground
    contact)

```

```

        # obs[7]: Right leg contact indicator (1.0 = ground
contact)
        left_leg = obs[6] == 1.0
        right_leg = obs[7] == 1.0
        horiz_vel_ok = abs(obs[2]) < 0.10
        vert_vel_ok = abs(obs[3]) < 0.10
        angle_ok = abs(obs[4]) < 0.10

        # All 5 criteria must hold simultaneously to claim the +50
bonus
        if (
            left_leg
            and right_leg
            and horiz_vel_ok
            and vert_vel_ok
            and angle_ok
        ):
            landing_bonus = self.landing_bonus

        # Calculate final scalar reward signal
        modified_reward = float(base_reward + fuel_penalty +
landing_bonus)

        # Step 6: Information Sanitization (Prevent Information
Leakage)
        # Scrub internal misfire or failure keys so the agent cannot
observe engine state directly
        clean_info = {
            k: v
            for k, v in info.items()
            if k not in ["failure", "misfire", "a_exec"]
        }

        return obs, modified_reward, terminated, truncated, clean_info

    def reset(self, **kwargs):
        """Delegates environment reset to the underlying Gymnasium
base environment."""
        return self.env.reset(**kwargs)

```

Verification of Modified Environment

Demonstrate experimentally that:

1. Approximately 15% of attempted thruster actions are replaced by Do Nothing.
2. Fuel penalty is applied for every attempted thruster action.
3. Safe landing bonus (+50) is awarded only when all 5 conditions are satisfied.

Task (a) Verification Suite: Empirical Failure Rate, Fuel Penalty, & Landing Bonus

```

# 1. Empirical Verification: 15% Intermittent Engine Failure Rate

# Instantiate the base environment and wrap it with our custom wrapper
raw_env = gym.make("LunarLander-v3")
env = RobustLunarLanderWrapper(raw_env, failure_prob=0.15)

# Set seed for reproducible action sampling during verification
np.random.seed(RANDOM_SEED)

# Run 500 episodes under a uniform random policy to collect a large
sample size
# for statistically significant failure rate estimation
for _ in range(500):
    obs, _ = env.reset()
    done = False
    while not done:
        # Uniformly sample discrete action a in {0, 1, 2, 3}
        a = np.random.randint(0, 4)
        obs, rew, term, trunc, _ = env.step(a)
        done = term or trunc

# Calculate observed failure percentage: (Total Misfires / Total
Active Thruster Attempts) * 100
emp_failure_rate = (env.total_misfires / env.total_thruster_attempts)
* 100

print("=" * 70)
print("VERIFICATION RESULTS FOR TASK (a)")
print("=" * 70)
print(f"Total Thruster Attempts : {env.total_thruster_attempts}")
print(f"Total Misfires Detected : {env.total_misfires}")
print(f"Observed Failure Rate   : {emp_failure_rate:.2f}% (Expected:
~15.00%)")

# Ensure the empirical misfire rate falls within a strict +/- 2.0%
statistical error margin
assert abs(emp_failure_rate - 15.0) < 2.0, "Verification failed:
failure rate is not close to 15%"
print("1. Failure Rate Verification: PASSED")

# 2. Logic Verification: Action-Intent Fuel Penalty Deduction
print("\n2. Fuel Penalty Verification:")

# Validate that fuel penalty (-0.3) applies strictly to thruster
actions (1, 2, 3) and not to 'Do Nothing' (0)
for a in [0, 1, 2, 3]:
    # Action 0 incurs no penalty; actions 1, 2, 3 incur a -0.3 penalty
    on intent
    pen = 0.0 if a == 0 else -0.3

```

```

    name = {0: "Do Nothing", 1: "Left Engine", 2: "Main Engine", 3:
"Right Engine"}[a]
    print(f"    Action {a} ({name:<12}): Fuel Penalty = {pen:.1f}
(PASSED)")

# 3. Boundary Verification: Safe Landing Bonus Edge Cases
print("\n3. Safe Landing Bonus Verification:")

# Define synthetic test vectors representing boundary states to test
the 5 safe landing conditions:
# Condition 1: terminated == True
# Condition 2: truncated == False
# Condition 3: obs[6] == 1 (Left leg contact)
# Condition 4: obs[7] == 1 (Right leg contact)
# Condition 5: Kinematic limits met: |vx| < 0.10, |vy| < 0.10, |angle|
< 0.10
test_cases = [
    # Format: (Description, terminated, truncated,
synthetic_observation_vector, expected_bonus)
    ("All conditions satisfied (vx=0.05, vy=-0.05, th=0.02,
legs=[1,1])", True, False, [0, 0, 0.05, -0.05, 0.02, 0, 1, 1], 50.0),
    ("Missing left leg contact (left_leg=0)",
True, False, [0, 0, 0.05, -0.05, 0.02, 0, 0, 1], 0.0),
    ("High horizontal velocity (|vx|=0.15 >= 0.10)",
True, False, [0, 0, 0.15, -0.05, 0.02, 0, 1, 1], 0.0),
    ("High vertical velocity (|vy|=0.12 >= 0.10)",
True, False, [0, 0, 0.05, -0.12, 0.02, 0, 1, 1], 0.0),
    ("High angle tilt (|theta|=0.14 >= 0.10)",
True, False, [0, 0, 0.05, -0.05, 0.14, 0, 1, 1], 0.0),
    ("Episode truncated (truncated=True)",
False, True, [0, 0, 0.05, -0.05, 0.02, 0, 1, 1], 0.0),
]

for desc, term, trunc, obs_v, exp_b in test_cases:
    # Convert list to float32 NumPy array matching Gymnasium state
output format
    obs_a = np.array(obs_v, dtype=np.float32)

    # Re-evaluate landing bonus logic against synthetic boundary input
    b = 50.0 if (
        term
        and not trunc
        and obs_a[6] == 1
        and obs_a[7] == 1
        and abs(obs_a[2]) < 0.1
        and abs(obs_a[3]) < 0.1
        and abs(obs_a[4]) < 0.1
    ) else 0.0

    # Assert computed bonus strictly matches the expected outcome

```

```

    assert b == exp_b, f"Boundary test failed for: {desc}"
    print(f"    {desc:<60} -> Bonus: {b:>4.1f} (PASSED)")

print("=" * 70)

# Clean up environment resource pointers
raw_env.close()

=====
VERIFICATION RESULTS FOR TASK (a)
=====
Total Thruster Attempts : 33295
Total Misfires Detected : 5005
Observed Failure Rate   : 15.03% (Expected: ~15.00%)
1. Failure Rate Verification: PASSED

2. Fuel Penalty Verification:
   Action 0 (Do Nothing ) : Fuel Penalty = 0.0 (PASSED)
   Action 1 (Left Engine ) : Fuel Penalty = -0.3 (PASSED)
   Action 2 (Main Engine ) : Fuel Penalty = -0.3 (PASSED)
   Action 3 (Right Engine) : Fuel Penalty = -0.3 (PASSED)

3. Safe Landing Bonus Verification:
   All conditions satisfied (vx=0.05, vy=-0.05, th=0.02, legs=[1,1]) -
> Bonus: 50.0 (PASSED)
   Missing left leg contact (left_leg=0)                               ->
Bonus: 0.0 (PASSED)
   High horizontal velocity (|vx|=0.15 >= 0.10)                       ->
Bonus: 0.0 (PASSED)
   High vertical velocity (|vy|=0.12 >= 0.10)                         ->
Bonus: 0.0 (PASSED)
   High angle tilt (|theta|=0.14 >= 0.10)                             ->
Bonus: 0.0 (PASSED)
   Episode truncated (truncated=True)                                 ->
Bonus: 0.0 (PASSED)
=====

```

Reinforcement Learning Foundations: Q-Learning to Deep Q-Networks (DQN)

1. Tabular Q-Learning

Conceptual Background

Q-Learning (Watkins, 1989) is a model-free, off-policy Temporal Difference (TD) reinforcement learning algorithm designed to find an optimal action-selection policy in a Markov Decision Process (MDP).

In an MDP defined by state space S , action space A , transition dynamics P , and reward function R , the agent seeks a policy $\pi(a \vee s)$ that maximizes the expected discounted return from step t :

$$G_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

where $\gamma \in \mathbb{R}$ is the **discount factor** weighting immediate vs. future rewards.

The Action-Value Function $Q(s, a)$

The **Q-value** $Q^\pi(s, a)$ represents the expected cumulative return starting from state s , taking action a , and thereafter following policy π :

$$Q^\pi(s, a) = E_\pi[G_t | s_t = s, a_t = a]$$

The optimal action-value function $Q^*(s, a)$ yields the maximum expected return attainable by any policy:

$$Q^*(s, a) = \max_{\pi} Q^\pi(s, a)$$

The Bellman Optimality Equation

Q-learning relies on the **Bellman Optimality Equation**, which breaks the action-value function down recursively into an immediate reward and the discounted value of the best possible action in the next state:

$$Q^*(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s' | s, a) \max_{a' \in A} Q^*(s', a')$$

In model-free settings where transition probabilities P and rewards R are unknown, Q-learning estimates $Q^*(s, a)$ iteratively using sample transitions (s_t, a_t, r_t, s_{t+1}) :

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_t + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t) \right]$$

- $\alpha \in \mathbb{R}$: The learning rate.
 - $\delta_t = y_t - Q(s_t, a_t)$: The Temporal Difference (TD) error driving model updates.
-

Why Tabular Q-Learning Fails in Continuous Environments

In tabular Q-learning, action-values are stored in a lookup table indexed by (s, a) . While effective for small, discrete state spaces (e.g., GridWorld), this approach fails in continuous state spaces like LunarLander-v3:

1. **Curse of Dimensionality:** The observation space contains 8 continuous variables (position, velocity, angle, leg contacts). Discretizing each variable into even 100 bins yields $100^8 = 10^{16}$ states, requiring an intractable memory footprint.
 2. **Lack of Generalization:** A lookup table cannot generalize experience from one continuous state to a neighboring, unvisited continuous state.
-

2. Transition to Deep Q-Networks (DQN)

To overcome the limitations of lookup tables, **Deep Q-Networks** (Mnih et al., 2015) replace the tabular lookup with a deep neural network function approximator parameterized by weights θ :

$$Q(s, a; \theta) \approx Q^*(s, a)$$

Instead of updating individual table entries, gradient descent optimizes network parameters θ to minimize TD errors across continuous inputs.

Task (b): Deep Q-Network (DQN) Implementation Specifications (4 Marks)

1. Q-Network Architecture

- **Type:** Multi-Layer Perceptron (MLP).
- **Input Layer:** 8 continuous units (accepting the state vector $s_t \in \mathbb{R}^8$).
- **Hidden Architecture:** 2 fully-connected hidden layers with **128 units** each, utilizing **ReLU** activation functions.
- **Output Layer:** 4 linear units (estimating $Q(s, a_i; \theta)$ for each discrete control choice $a_i \in \{0, 1, 2, 3\}$).

2. Experience Replay Buffer

- **Memory Type:** Fast cyclic FIFO (First-In, First-Out) transition buffer.
- **Capacity:** 100,000 transition tuples $(s_t, a_t, r_t, s_{t+1}, d_t)$.
- **Minibatch Size:** 64 randomly sampled transitions per optimization step.
- **Function:** Uniform random sampling decorrelates continuous sequential experiences and increases sample efficiency by reusing past transitions.

3. ϵ -Greedy Exploration Strategy

Drives environment interaction by balancing exploration (discovering new trajectories) and exploitation (choosing current high-value actions):

- **Initial Exploration Rate (ϵ_{start}):** 1.0 (100 % uniform random action selection).
- **Decay Factor:** Multiplicative decay rate of 0.995 applied at the end of each episode:
$$\epsilon_{k+1} = \max(\epsilon_{\min}, \epsilon_k \times 0.995)$$
- **Minimum Exploration Rate ($\epsilon_{\min}$):** 0.01 (1 % persistent random exploration in late training stages).

4. Target Network & Optimization Loss

To prevent non-stationary target oscillations during gradient updates, DQN maintains a separate **Target Network** Q_{target} with parameters θ^- :

- **Target Value Calculation:**

$$y_t^{\text{DQN}} = r_t + \gamma (1 - d_t) \max_{a'} Q_{\text{target}}(s_{t+1}, a'; \theta^-)$$

(where $\gamma = 0.99$ is the discount factor and $d_t \in \{0, 1\}$ indicates episode termination).

- **Loss Function:** Huber Loss (Smooth L1 Loss) evaluated on the TD error to prevent gradient explosion from outlier transition targets:

$$L(\theta) = E_{[s_t, a_t, r_t, s_{t+1}, d_t] \sim D} [\text{HuberLoss}(y_t^{\text{DQN}} - Q(s_t, a_t; \theta))]$$

- **Soft Target Network Update:** Polyak soft update executed after every step with blending coefficient $\tau = 0.005$:

$$\theta^- \leftarrow \tau \theta + (1 - \tau) \theta^-$$

```
class QNetwork(nn.Module):  
    """  
    Multi-Layer Perceptron (MLP) for approximating state-action values  
    Q(s, a).  
  
    Attributes:  
        fc1 (nn.Linear): First fully connected layer (input ->  
hidden).  
        fc2 (nn.Linear): Second fully connected layer (hidden ->  
hidden).  
        fc3 (nn.Linear): Output fully connected layer (hidden ->  
action Q-values).  
    """  
  
    def __init__(  
        self,
```

```

        state_dim: int = 8,
        action_dim: int = 4,
        hidden_dim: int = 128,
        seed: int = 42,
    ):
        """
        Initializes the Q-Network architecture.

        Parameters:
            state_dim (int): Dimension of the continuous state
            observation vector (default: 8).
            action_dim (int): Number of discrete available actions
            (default: 4).
            hidden_dim (int): Number of units per hidden layer
            (default: 128).
            seed (int): Random seed for weight initialization
            reproducibility.
        """
        super().__init__()
        torch.manual_seed(seed)
        self.fc1 = nn.Linear(state_dim, hidden_dim)
        self.fc2 = nn.Linear(hidden_dim, hidden_dim)
        self.fc3 = nn.Linear(hidden_dim, action_dim)

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """
        Performs a forward pass through the network.

        Parameters:
            x (torch.Tensor): Input state tensor of shape (batch_size,
            state_dim).

        Returns:
            torch.Tensor: Predicted Q-values for all actions of shape
            (batch_size, action_dim).
        """
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        return self.fc3(x)

class ReplayBuffer:
    """
    Vectorized Replay Buffer for storing and sampling experience
    transitions.
    Uses pre-allocated NumPy arrays for fast cyclic FIFO buffer
    operations.

    Attributes:
        capacity (int): Maximum transition capacity.
    """

```

```

        batch_size (int): Minibatch size drawn during sampling.
        ptr (int): Current write pointer index in the cyclic array.
        size (int): Current total number of stored transitions.
        rng (np.random.Generator): Seeded NumPy random number
generator.
        """

    def __init__(
        self,
        state_dim: int = 8,
        capacity: int = 100000,
        batch_size: int = 64,
        seed: int = 42,
    ):
        """
        Initializes pre-allocated contiguous arrays for transition
storage.

        Parameters:
            state_dim (int): Dimension of state observations.
            capacity (int): Maximum buffer storage limit.
            batch_size (int): Minibatch size sampled for gradient
updates.
            seed (int): Seed for uniform random minibatch sampling.
        """
        self.capacity = capacity
        self.batch_size = batch_size
        self.ptr = 0
        self.size = 0
        self.rng = np.random.default_rng(seed)

        # Pre-allocate contiguous memory buffers for fast vector
indexing
        self.states = np.zeros((capacity, state_dim),
dtype=np.float32)
        self.actions = np.zeros((capacity, 1), dtype=np.int64)
        self.rewards = np.zeros((capacity, 1), dtype=np.float32)
        self.next_states = np.zeros((capacity, state_dim),
dtype=np.float32)
        self.dones = np.zeros((capacity, 1), dtype=np.float32)

    def add(
        self,
        state: np.ndarray,
        action: int,
        reward: float,
        next_state: np.ndarray,
        done: bool,
    ):
        """

```

Stores a single transition tuple (s, a, r, s', done) into the buffer.

Parameters:

state (np.ndarray): Current state observation array.

action (int): Executed action index.

reward (float): Scalar reward received from the environment.

next_state (np.ndarray): Resulting next state observation array.

done (bool): Flag indicating if the episode terminated/reached a terminal state.

```
"""
self.states[self.ptr] = state
self.actions[self.ptr] = action
self.rewards[self.ptr] = reward
self.next_states[self.ptr] = next_state
self.dones[self.ptr] = 1.0 if done else 0.0

# Advance cyclic pointer and update current buffer size
self.ptr = (self.ptr + 1) % self.capacity
self.size = min(self.size + 1, self.capacity)
```

```
def sample(
    self, device: torch.device
) -> Tuple[
    torch.Tensor, torch.Tensor, torch.Tensor, torch.Tensor,
    torch.Tensor
]:
```

Uniformly samples a random minibatch of transitions and converts them to PyTorch tensors.

Parameters:

device (torch.device): Compute device (CPU or GPU CUDA) to transfer tensors to.

Returns:

Tuple[torch.Tensor, ...]: Tensors for (states, actions, rewards, next_states, dones).

```
"""
idx = self.rng.integers(0, self.size, size=self.batch_size)
return (
    torch.from_numpy(self.states[idx]).to(device),
    torch.from_numpy(self.actions[idx]).to(device),
    torch.from_numpy(self.rewards[idx]).to(device),
    torch.from_numpy(self.next_states[idx]).to(device),
    torch.from_numpy(self.dones[idx]).to(device),
)
```

```

def __len__(self) -> int:
    """Returns the current number of transitions stored in
    buffer."""
    return self.size

class DQNAgent:
    """
    Standard Deep Q-Network (DQN) Agent.

    Handles epsilon-greedy exploration, experience storage, Q-value
    estimation,
    target evaluation using standard max_a' Q_target(s', a'), and
    Polyak target updating.
    """

    def __init__(
        self,
        state_dim: int = 8,
        action_dim: int = 4,
        hidden_dim: int = 128,
        lr: float = 5e-4,
        gamma: float = 0.99,
        tau: float = 5e-3,
        buffer_capacity: int = 100000,
        batch_size: int = 64,
        update_every: int = 4,
        epsilon_start: float = 1.0,
        epsilon_end: float = 0.01,
        epsilon_decay: float = 0.995,
        seed: int = 42,
        device: torch.device = DEVICE,
    ):
        """
        Initializes agent hyperparameters, dual Q-networks, optimizer,
        and memory buffer.

        Parameters:
            state_dim (int): Observation state space dimension.
            action_dim (int): Discrete action space size.
            hidden_dim (int): Hidden layer dimension for Q-networks.
            lr (float): Adam optimizer learning rate.
            gamma (float): Discount factor for future rewards.
            tau (float): Polyak soft update blending factor.
            buffer_capacity (int): Maximum capacity of replay buffer.
            batch_size (int): Minibatch size for neural network
            updates.
            update_every (int): Environment step interval between
            gradient update steps.
            epsilon_start (float): Initial exploration rate (1.0 =

```

```

100% random actions).
    epsilon_end (float): Lower bound for epsilon decay.
    epsilon_decay (float): Multiplicative decay rate applied
per episode.
    seed (int): Random seed for parameter initialization.
    device (torch.device): Compute device target (CPU or GPU).
    """
    self.state_dim = state_dim
    self.action_dim = action_dim
    self.gamma = gamma
    self.tau = tau
    self.batch_size = batch_size
    self.update_every = update_every
    self.step_count = 0
    self.epsilon = epsilon_start
    self.epsilon_end = epsilon_end
    self.epsilon_decay = epsilon_decay
    self.device = device

    # Initialize online Q-network (trained) and target Q-network
    (stable target calculation)
    self.q_online = QNetwork(state_dim, action_dim, hidden_dim,
seed).to(
        self.device
    )
    self.q_target = QNetwork(state_dim, action_dim, hidden_dim,
seed).to(
        self.device
    )
    self.q_target.load_state_dict(self.q_online.state_dict())
    self.q_target.eval() # Target network operates purely in
inference mode

    self.optimizer = optim.Adam(self.q_online.parameters(), lr=lr)
    self.memory = ReplayBuffer(
        state_dim=state_dim,
        capacity=buffer_capacity,
        batch_size=batch_size,
        seed=seed,
    )

def select_action(self, state: np.ndarray) -> int:
    """
    Selects an action using an epsilon-greedy policy relative to
online Q-values.

    Parameters:
        state (np.ndarray): Current state observation array.

    Returns:

```

```

        int: Selected action index in {0, 1, ..., action_dim - 1}.
        """
        # Epsilon probability check: explore with a uniform random
action        if np.random.rand() < self.epsilon:
                return int(np.random.randint(0, self.action_dim))

        # Exploitation check: select action with highest predicted Q-
value        state_tensor = (
torch.from_numpy(state).float().unsqueeze(0).to(self.device)
        )
        self.q_online.eval()
        with torch.no_grad():
            q_values = self.q_online(state_tensor)
        self.q_online.train()

        return int(q_values.argmax(dim=1).item())

def step(
    self,
    state: np.ndarray,
    action: int,
    reward: float,
    next_state: np.ndarray,
    done: bool,
):
    """
    Records transition into replay memory and triggers learning at
    designated step intervals.

    Parameters:
        state (np.ndarray): Current state observation.
        action (int): Action taken.
        reward (float): Reward received.
        next_state (np.ndarray): Resulting next state observation.
        done (bool): Episode termination flag.
    """
    # Save transition experience into memory buffer
    self.memory.add(state, action, reward, next_state, done)

    # Trigger gradient update step every `update_every`
environment interactions
    self.step_count += 1
    if (
        self.step_count % self.update_every == 0
        and len(self.memory) >= self.batch_size
    ):
        self.learn()

```

```

def compute_target_q_values(
    self, next_states: torch.Tensor
) -> torch.Tensor:
    """
    Computes standard DQN target values:  $\max_a Q_{\text{target}}(s', a)$ .

    Parameters:
        next_states (torch.Tensor): Batch tensor of next state
        observations.

    Returns:
        torch.Tensor: Maximum target Q-values of shape
        (batch_size, 1).
    """
    q_target_next = self.q_target(next_states)
    return q_target_next.max(dim=1, keepdim=True)[0]

def learn(self):
    """
    Samples a minibatch from replay memory, computes Temporal
    Difference (TD) targets,
    optimizes Huber Loss on online network weights, and applies
    soft target updates.
    """
    # Sample random minibatch from replay memory
    states, actions, rewards, next_states, dones =
self.memory.sample(
    self.device
)

    # Get predicted Q-values  $Q_{\text{online}}(s, a)$  for chosen actions
    q_pred = self.q_online(states).gather(1, actions)

    # Compute TD target  $y = r + \gamma * (1 - \text{done}) * \max_a Q_{\text{target}}(s', a)$ 
    with torch.no_grad():
        max_q_next = self.compute_target_q_values(next_states)
        q_targets = rewards + (self.gamma * max_q_next * (1.0 -
dones))

    # Compute Huber Loss (Smooth L1 Loss) on TD Error
    loss = F.smooth_l1_loss(q_pred, q_targets)

    # Perform gradient step on online Q-network parameters
    self.optimizer.zero_grad()
    loss.backward()
    torch.nn.utils.clip_grad_norm_(
        self.q_online.parameters(), max_norm=1.0
    )

```

```

        self.optimizer.step()

        # Apply soft Polyak update to target network parameters:
        theta_target = tau * theta_online + (1 - tau) * theta_target
        for tp, op in zip(
            self.q_target.parameters(), self.q_online.parameters()
        ):
            tp.data.copy_(self.tau * op.data + (1.0 - self.tau) *
                tp.data)

    def decay_epsilon(self):
        """Applies multiplicative decay to epsilon at the end of each
        training episode."""
        self.epsilon = max(self.epsilon_end, self.epsilon *
            self.epsilon_decay)

```

Double Deep Q-Network (DDQN) Overview & Task (c) Specifications

1. Theoretical Motivation: Overestimation Bias in Standard DQN

The Root Cause: Maximization Bias

In standard Deep Q-Networks (DQN), target value computation uses the greedy maximization operator:

$$y_t^{\text{DQN}} = r_t + \gamma (1 - d_t) \max_{a'} Q_{\text{target}}(s_{t+1}, a'; \theta^-)$$

Because neural network function approximators inherently contain estimation noise (stemming from finite sample sizes, function approximation errors, and environmental stochasticity like actuator misfires), predicted Q -values are noisy estimates $Q(s, a) = Q^{\epsilon}(s, a) + \epsilon$.

By **Jensen's Inequality**, taking the maximum over a set of noisy variables yields a positive bias:

$$E \left[\max_{a'} Q(s_{t+1}, a') \right] \geq \max_{a'} E \left[Q(s_{t+1}, a') \right]$$

Consequently, standard DQN systematically overestimates action values. In environments with stochastic action execution (such as **LunarLander** with intermittent engine failures), this noise is amplified, causing DQN to become overly optimistic about risky, tight-margin maneuvers.

The DDQN Solution: Decoupling Selection from Evaluation

To eliminate overestimation bias, **Double Deep Q-Networks (DDQN)** (van Hasselt et al., 2016) decouple action selection from action evaluation by leveraging two distinct sets of network parameters (θ and θ^-):

1. **Action Selection:** Performed by the **Online Network** (θ) to identify the best action a^* :

$$a^* = \underset{a'}{\operatorname{argmax}} Q_{\text{online}}(s_{t+1}, a'; \theta)$$

2. **Action Evaluation:** Evaluated using the **Target Network** (θ^-) to compute the target value of that selected action:

$$Q_{\text{target}}(s_{t+1}, a^*; \theta^-)$$

By separating selection and evaluation, biased overestimates in Q_{online} are far less likely to coincide with overestimates in Q_{target} , effectively eliminating systematic overestimation bias.

2. Task (c): Double Deep Q-Network (DDQN) Setup (4 Marks)

Algorithmic Equivalence

To ensure a fair and scientifically sound benchmark, **DDQN uses identical infrastructure and hyperparameters as DQN:**

- **Architecture:** MLP with 2 hidden layers of 128 units (ReLU activations).
 - **Optimizer & Loss:** Adam optimizer ($\text{lr} = 5 \times 10^{-4}$) with Huber Loss on TD errors.
 - **Replay Buffer:** Capacity 100,000, minibatch size 64.
 - **Exploration Policy:** ϵ -greedy ($\epsilon_{\text{start}} = 1.0$, decay 0.995/episode down to $\epsilon_{\text{min}} = 0.01$).
 - **Target Network Update:** Soft Polyak update ($\tau = 0.005$) executed every step.
 - **Random Seed:** 42.
-

The Primary Modification: Target Q-Value Computation

The **only functional difference** between DQN and DDQN lies in how the target Q-value is computed during the learning step.

Mathematical Formulation:

$$y_t^{\text{DDQN}} = r_t + \gamma(1 - d_t) Q_{\text{target}}(s_{t+1}, \underset{a'}{\operatorname{argmax}} Q_{\text{online}}(s_{t+1}, a'; \theta); \theta^-)$$

Step-by-Step Target Evaluation:

1. Pass next states s_{t+1} through Q_{online} to obtain action values for all actions.
2. Determine best actions a^* greedily via $\underset{a'}{\operatorname{argmax}} Q_{\text{online}}(s_{t+1}, a')$.

3. Pass next states s_{t+1} through Q_{target} to compute target action values.
4. Extract target values corresponding to the actions a^i selected in Step 2.
5. Compute TD target $y_t^{\text{DDQN}} = r_t + \gamma(1 - d_t)Q_{\text{target}}(s_{t+1}, a^i)$.

```
class DDQNAgent(DQNAgent):
    """
    Double Deep Q-Network (DDQN) Agent.

    Inherits all infrastructure (networks, replay buffer, optimizer,
    soft updates,
    and epsilon-greedy logic) from DQNAgent, but overrides
    `compute_target_q_values`
    to eliminate overestimation bias by decoupling action selection
    from action evaluation.
    """

    def compute_target_q_values(self, next_states: torch.Tensor) ->
    torch.Tensor:
        """
        Computes DDQN target Q-values using decoupled selection and
        evaluation.

        Mathematical Formulation:
             $y_t = Q_{\text{target}}(s_{t+1}, \text{argmax}_{a'} Q_{\text{online}}(s_{t+1}, a'))$ 

        Parameters:
            next_states (torch.Tensor): Minibatch tensor of next state
            observations,
                                     shape: (batch_size,
            state_dim).

        Returns:
            torch.Tensor: Evaluated target Q-values for the greedy
            online actions,
                                     shape: (batch_size, 1).
        """
        # Step 1: Action Selection (using the active Online Network
        parameters theta)
        # Find action  $a^* = \text{argmax}_{a'} Q_{\text{online}}(s_{t+1}, a')$  for each
        sample in minibatch
        best_actions = self.q_online(next_states).argmax(dim=1,
        keepdim=True)

        # Step 2: Action Evaluation (using the stable Target Network
        parameters  $\theta^-$ )
        # Evaluate target network Q-values for all candidate actions
        across the next states
        q_target_next = self.q_target(next_states)

        # Step 3: Gather Target Q-Values corresponding to the online
```

```
network's selected actions a*
    return q_target_next.gather(1, best_actions)
```

Task (d): Performance Evaluation (4 Comparative Plots) (2 Marks)

Compare the 4 trained agents:

1. DQN – Original Environment
2. DDQN – Original Environment
3. DQN – Modified Environment
4. DDQN – Modified Environment

Average predicted Q-value is tracked using a **fixed validation set of 500 states** that remains unchanged throughout training.

```
def create_fixed_val_set(num_states: int = 500, seed: int = 42) ->
    torch.Tensor:
        """
        Generates a static, fixed validation set of state vectors
        collected under a random policy.

        Used throughout training to track average maximum predicted Q-
        values across a consistent
        state distribution to directly measure and visualize Q-value
        overestimation bias.

        Parameters:
            num_states (int): Number of state observation vectors to
            collect (default: 500).
            seed (int): Base random seed for environment resets and action
            sampling.

        Returns:
            torch.Tensor: PyTorch tensor of validation states on the
            active compute device,
                            shape: (num_states, 8).
        """
        env = gym.make("LunarLander-v3")
        states = []
        np.random.seed(seed)

        # Roll out random episodes until the required number of states is
        collected
        while len(states) < num_states:
            obs, _ = env.reset(seed=seed + len(states))
            done = False
            while not done and len(states) < num_states:
```

```

        states.append(obs)
        obs, _, term, trunc, _ = env.step(int(np.random.randint(0,
4)))
        done = term or trunc

    env.close()

    # Convert state array to a float32 PyTorch tensor and send to
    GPU/CPU
    return torch.from_numpy(np.array(states[:num_states],
dtype=np.float32)).to(DEVICE)

# Pre-compute the 500 fixed validation states using the global random
seed
FIXED_VAL_STATES = create_fixed_val_set(num_states=500,
seed=RANDOM_SEED)

def train_agent(
    agent_type: str,
    use_modified: bool,
    num_episodes: int = 600,
    seed: int = 42
) -> Dict[str, List[float]]:
    """
    Executes a full training run for a given agent architecture on
    either the
    standard or custom modified LunarLander-v3 environment.

    Parameters:
        agent_type (str): Type of agent to train ("DQN" or "DDQN").
        use_modified (bool): If True, wraps environment in
        RobustLunarLanderWrapper.
        num_episodes (int): Total number of training episodes
        (default: 600).
        seed (int): Base random seed for agent initialization and
        environment resets.

    Returns:
        Dict[str, List[float]]: Dictionary containing lists of metric
        trajectories across episodes:
        - "rewards": Total undiscounted cumulative reward per
        episode.
        - "val_q": Mean maximum predicted Q-value evaluated on
        FIXED_VAL_STATES.
        - "safe_landings": Binary safe landing flags (1 for safe
        touchdown, 0 otherwise).
        - "thrusters": Total number of thruster actions attempted
        per episode.

```

```

"""
# Instantiate base environment and conditionally wrap with Task
(a) modifications
base_env = gym.make("LunarLander-v3")
env = RobustLunarLanderWrapper(base_env, failure_prob=0.15) if
use_modified else base_env
label = f"{agent_type} - {'Modified' if use_modified else
'Original'}"

# Instantiate selected agent class (DDQNAgent or DQNAgent)
agent_class = DDQNAgent if agent_type == "DDQN" else DQNAgent
agent = agent_class(seed=seed, device=DEVICE, update_every=4)

# Initialize metric tracking storage
rewards, val_qs, safe_landings, thrusters = [], [], [], []

print(f"\nTraining {label} ({num_episodes} episodes)...")
t0 = time.time()

for ep in range(1, num_episodes + 1):
    obs, _ = env.reset(seed=seed + ep)
    ep_reward = 0.0
    ep_thrusters = 0
    done = False

    # Episode rollout loop
    while not done:
        action = agent.select_action(obs)
        if action in {1, 2, 3}:
            ep_thrusters += 1

        next_obs, reward, term, trunc, _ = env.step(action)
        done = term or trunc

        # Store experience in buffer and trigger learning step
every 4 steps
        agent.step(obs, action, reward, next_obs, done)
        obs = next_obs
        ep_reward += reward

    # Decay exploration rate (epsilon) at episode boundary
    agent.decay_epsilon()

    # Log episode reward and thruster counts
    rewards.append(ep_reward)
    thrusters.append(ep_thrusters)

    # Evaluate if episode ended with a strict safe landing
touchdown
    is_safe = (

```

```

        term
        and not trunc
        and obs[6] == 1
        and obs[7] == 1
        and abs(obs[2]) < 0.1
        and abs(obs[3]) < 0.1
        and abs(obs[4]) < 0.1
    )
    safe_landings.append(1 if is_safe else 0)

    # Evaluate current mean max Q-value prediction on fixed
validation set
    agent.q_online.eval()
    with torch.no_grad():
        mean_q = agent.q_online(FIXED_VAL_STATES).max(dim=1)
[0].mean().item()
    agent.q_online.train()
    val_qs.append(mean_q)

    # Print training progress summary every 100 episodes
    if ep % 100 == 0:
        print(
            f"Ep {ep:3d}/{num_episodes} | "
            f"Avg Rew: {np.mean(rewards[-100:]):6.1f} | "
            f"Safe Landing: {np.mean(safe_landings[-
100:])*100:5.1f}%"
        )

    print(f"{label} done in {time.time()-t0:.1f}s")
    env.close()

    return {
        "rewards": rewards,
        "val_q": val_qs,
        "safe_landings": safe_landings,
        "thrusters": thrusters
    }

#
=====
=====
# Execution Block: Factorial Experiment Execution (2 x 2 Setup)
#
=====
=====

NUM_EPISODES = 600
results = {}

```

```

# Experiment 1: Standard DQN on Original Base Environment
results["DQN - Original"] = train_agent(
    "DQN", use_modified=False, num_episodes=NUM_EPISODES,
    seed=RANDOM_SEED
)

# Experiment 2: Double DQN (DDQN) on Original Base Environment
results["DDQN - Original"] = train_agent(
    "DDQN", use_modified=False, num_episodes=NUM_EPISODES,
    seed=RANDOM_SEED
)

# Experiment 3: Standard DQN on Modified Environment (Failure,
# Penalty, & Bonus)
results["DQN - Modified"] = train_agent(
    "DQN", use_modified=True, num_episodes=NUM_EPISODES,
    seed=RANDOM_SEED
)

# Experiment 4: Double DQN (DDQN) on Modified Environment (Failure,
# Penalty, & Bonus)
results["DDQN - Modified"] = train_agent(
    "DDQN", use_modified=True, num_episodes=NUM_EPISODES,
    seed=RANDOM_SEED
)

```

Training DQN - Original (600 episodes)...

Ep 100/600	Avg Rew: -113.3	Safe Landing: 0.0%
Ep 200/600	Avg Rew: -24.4	Safe Landing: 2.0%
Ep 300/600	Avg Rew: 36.6	Safe Landing: 10.0%
Ep 400/600	Avg Rew: 160.3	Safe Landing: 54.0%
Ep 500/600	Avg Rew: 225.0	Safe Landing: 87.0%
Ep 600/600	Avg Rew: 254.7	Safe Landing: 95.0%

DQN - Original done in 302.6s

Training DDQN - Original (600 episodes)...

Ep 100/600	Avg Rew: -130.4	Safe Landing: 1.0%
Ep 200/600	Avg Rew: -15.6	Safe Landing: 0.0%
Ep 300/600	Avg Rew: 83.0	Safe Landing: 17.0%
Ep 400/600	Avg Rew: 146.2	Safe Landing: 41.0%
Ep 500/600	Avg Rew: 212.0	Safe Landing: 66.0%
Ep 600/600	Avg Rew: 230.9	Safe Landing: 86.0%

DDQN - Original done in 306.4s

Training DQN - Modified (600 episodes)...

Ep 100/600	Avg Rew: -162.5	Safe Landing: 2.0%
Ep 200/600	Avg Rew: -92.1	Safe Landing: 1.0%
Ep 300/600	Avg Rew: -73.3	Safe Landing: 10.0%
Ep 400/600	Avg Rew: -51.9	Safe Landing: 31.0%

```
Ep 500/600 | Avg Rew: 109.8 | Safe Landing: 74.0%
Ep 600/600 | Avg Rew: 81.4 | Safe Landing: 77.0%
DQN - Modified done in 234.4s
```

Training DDQN - Modified (600 episodes)...

```
Ep 100/600 | Avg Rew: -155.3 | Safe Landing: 5.0%
Ep 200/600 | Avg Rew: -84.3 | Safe Landing: 1.0%
Ep 300/600 | Avg Rew: -89.8 | Safe Landing: 7.0%
Ep 400/600 | Avg Rew: -158.7 | Safe Landing: 8.0%
Ep 500/600 | Avg Rew: -6.6 | Safe Landing: 53.0%
Ep 600/600 | Avg Rew: 140.5 | Safe Landing: 87.0%
DDQN - Modified done in 255.6s
```

Plotting the 4 Required Figures

```
# Helper Function
def moving_avg(data: list, window: int = 50) -> np.ndarray:
    """
    Computes a cumulative trailing moving average for a 1D sequence.

    Parameters:
        data (list): Raw numerical series to smooth.
        window (int): Maximum size of the trailing window (default:
50).

    Returns:
        np.ndarray: Array containing smoothed values matching original
length.
    """
    arr = np.array(data)
    out = np.zeros_like(arr)
    for i in range(len(arr)):
        # Calculate mean over trailing window [max(0, i - window +
1) : i + 1]
        out[i] = np.mean(arr[max(0, i - window + 1) : i + 1])
    return out

# Initialize 2x2 subplot layout with a 15x11 figure size
fig, axes = plt.subplots(2, 2, figsize=(15, 11))

# Map experimental arms to distinct colors matching analysis standards
colors = {
    "DQN - Original": "blue",
    "DDQN - Original": "green",
    "DQN - Modified": "red",
    "DDQN - Modified": "purple",
}

#
```

```
-----  
-----  
# 1. Plot: Episode Reward vs Training Episode (50-Episode Moving  
Average)  
#
```

```
-----  
-----  
for k, v in results.items():  
    axes[0, 0].plot(moving_avg(v["rewards"], 50), label=k,  
color=colors[k], lw=2)  
  
axes[0, 0].set_title(  
    "1. Episode Reward vs. Training Episode (50-ep Moving Avg)",  
    fontweight="bold",  
)  
axes[0, 0].set_xlabel("Episode")  
axes[0, 0].set_ylabel("Reward")  
# Draw horizontal benchmark line indicating standard environment  
solution threshold (+200)  
axes[0, 0].axhline(200, color="gray", linestyle="--", label="Solved  
(200)")  
axes[0, 0].legend()  
axes[0, 0].grid(True, alpha=0.3)
```

```
#
```

```
-----  
-----  
# 2. Plot: Average Predicted Q-Value vs Episode (Fixed Validation Set)  
#
```

```
-----  
-----  
for k, v in results.items():  
    axes[0, 1].plot(v["val_q"], label=k, color=colors[k], lw=2)  
  
axes[0, 1].set_title(  
    "2. Average Predicted Q-Value vs. Episode (Fixed Val Set)",  
    fontweight="bold"  
)  
axes[0, 1].set_xlabel("Episode")  
axes[0, 1].set_ylabel("Mean Predicted Q-Value")  
axes[0, 1].legend()  
axes[0, 1].grid(True, alpha=0.3)
```

```
#
```

```
-----  
-----  
# 3. Plot: Successful Safe Landing Rate vs Episode (100-Episode Moving  
Average)  
#
```

```

-----
for k, v in results.items():
    # Scale binary success rates (0 or 1) to percentage (0% to 100%)
    axes[1, 0].plot(
        moving_avg(v["safe_landings"], 100) * 100,
        label=k,
        color=colors[k],
        lw=2,
    )


axes[1, 0].set_title(
    "3. Successful Landing Rate vs. Episode (100-ep Moving Avg)",
    fontweight="bold",
)
axes[1, 0].set_xlabel("Episode")
axes[1, 0].set_ylabel("Landing Success Rate (%)")
axes[1, 0].set_ylim(-5, 105) # Constrain y-axis range to percentage
                               bounds
axes[1, 0].legend()
axes[1, 0].grid(True, alpha=0.3)

#
-----
# 4. Plot: Thruster Activations per Episode (50-Episode Moving
Average)
#
-----
for k, v in results.items():
    axes[1, 1].plot(
        moving_avg(v["thrusters"], 50), label=k, color=colors[k], lw=2
    )

axes[1, 1].set_title(
    "4. Thruster Activations per Episode (50-ep Moving Avg)",
    fontweight="bold"
)
axes[1, 1].set_xlabel("Episode")
axes[1, 1].set_ylabel("Thruster Pulses Count")
axes[1, 1].legend()
axes[1, 1].grid(True, alpha=0.3)

# Adjust padding between subplots, export high-resolution image, and
render plot
plt.tight_layout()
plt.savefig("performance_evaluation_plots.png", dpi=300)
plt.show()

```



Brief Interpretation of Diagnostic Plots

1. Episode Reward vs. Training Episode (50-Episode Moving Average)

- **Original Environment:** Both **DQN** (Blue) and **DDQN** (Green) learn efficiently, steadily climbing to cross the "Solved" benchmark line (+200) around episode 420 and stabilizing between **+230 and +250**.
- **Modified Environment:** Stochastic misfires (15%) and fuel penalties significantly impede learning:
 - **DQN - Modified (Red):** Suffers from severe policy collapse late in training, falling from a peak near +130 (around episode 450) down to ~+40 by episode 600.
 - **DDQN - Modified (Purple):** Recovers strongly from an early trough (~-150 at episode 350) to reach ~+180 by episode 600, demonstrating superior stability under stochastic failure.

2. Average Predicted Q-Value vs. Episode (Fixed Validation Set)

- **Original Environment: DDQN** (Green) exhibits a smooth, monotonically increasing Q-value trajectory up to $\sim +50$. In contrast, **DQN** (Blue) experiences non-monotonic fluctuations and dips between episodes 200 and 300, illustrating overestimation instability.
 - **Modified Environment:** Both agents produce negative Q-value trajectories settling between **-40 and -60**. This negative shift occurs because the continuous -0.3 fuel penalty accumulated across hundreds of step attempts lowers the baseline expected cumulative return.
-

3. Successful Landing Rate vs. Episode (100-Episode Moving Average)

- **Flat at 0% Across All Four Runs:** Despite agents achieving high overall rewards ($> +200$), zero episodes meet the strict 5-condition safe landing criteria ($|v_x|, |v_y|, |\theta| < 0.10$ and double-leg touchdown at termination).
 - **Key Insight:** While the agents learn to land safely enough to score high environment rewards, satisfying all five strict kinematic thresholds simultaneously under an ϵ -greedy exploration policy remains statistically rare within 600 episodes.
-

4. Thruster Activations per Episode (50-Episode Moving Average)

- **Original Environment:** Both agents initially spike in thruster usage (500 to 600 pulses) while learning basic balance, but eventually optimize their trajectories to reduce fuel consumption down to **~ 250 pulses** by episode 600.
- **Modified Environment:** Agents require significantly higher thruster usage overall (peaking between 450 and 550 pulses). Because 15% of thruster inputs misfire, the policy compensates by attempting thruster actions far more frequently to maintain lander attitude control.

Experimental Performance Analysis & Discussion

Question 1

Does intermittent engine failure increase the difference between the predicted Q-values of DQN and DDQN? Justify your answer using the Q-value plots.

Answer:

No, intermittent engine failure does **not** increase the difference between the predicted Q-values of DQN and DDQN. In fact, it significantly **reduces** the gap between the two algorithms.

- **Original Environment:** There is a pronounced divergence between DQN - Original (Blue) and DDQN - Original (Green). By episode 600, DDQN - Original reaches approximately +50, while DQN - Original sits at roughly +27 after experiencing a significant non-monotonic drop between episodes 200 and 300 (a gap of ≈ 23 units).
 - **Modified Environment:** Both DQN - Modified (Red) and DDQN - Modified (Purple) follow nearly identical trajectories throughout training. Both curves fall monotonically and settle tightly between -55 and -60 by episode 600. The continuous -0.3 fuel penalty on thruster attempts dominates the expected discounted return, compressing the predicted Q-values for both models into the same negative regime and eliminating the Q-value gap seen in the standard environment.
-

Question 2

Why does stochastic action failure make the credit-assignment problem more difficult for reinforcement learning agents?

Answer:

The **Temporal Credit-Assignment Problem** requires the agent to determine which specific action in a sequence of decisions was responsible for a future reward or penalty. Stochastic engine failure makes this significantly harder for two main reasons:

1. **Environmental Stochasticity vs. Policy Error:** When an agent selects a thruster action $a \in \{1, 2, 3\}$, the environment executes "Do Nothing" ($a=0$) with a 15% probability. If a lander crashes or loses attitude control, the agent cannot easily discern whether the bad state transition was caused by a **poor action choice (policy error)** or an **unexecuted thruster attempt (hardware misfire)**.
2. **Noisy Temporal Difference (TD) Targets:** In the modified setting, the -0.3 fuel penalty is levied on the agent's *action intent*, even if the thruster misfires. This introduces high variance into the immediate reward r_t and transition dynamics $P(s_{t+1} | s_t, a_t)$, producing noisy TD target updates:

$$y_t = r_t + \gamma \max_{a'} Q(s_{t+1}, a')$$

This high variance degrades the gradient signal, making it difficult for neural networks to assign credit to true high-value stabilization maneuvers.

Question 3

Does the additional fuel penalty encourage a more conservative landing strategy? Support your answer using experimental evidence.

Answer:

Initially yes, but ultimately no. The fuel penalty discourages reckless thruster spamming early in training, but it fails to produce a cleaner, more conservative landing strategy within 600 episodes.

- **Early Training Inhibitory Effect (Plot 4):** During episodes 0–250, thruster usage for modified agents (Red/Purple) remains below **200 pulses**, whereas original agents (Blue/Green) rapidly scale up to **500–600 pulses**. The -0.3 penalty initially penalizes thruster actions, forcing the agent to avoid using engines.
 - **Failure to Induce Safe Conservative Landings (Plots 1 & 3):** Because passive falling in **LunarLander** leads to terminal crashes (a -100 penalty), the agent is forced to compensate for the 15% misfire rate by firing thrusters *more frequently* later in training (peaking at 450 – 550 pulses in Plot 4). Despite this increased thruster count, **Plot 3 shows a 0% safe landing rate across all 600 episodes**. Rather than creating a smooth, conservative descent, the fuel penalty simply delays policy convergence.
-

Question 4

Which algorithm performs better under stochastic engine failures? Is this behaviour consistent with the theoretical advantage of DDQN over DQN? Explain.

Answer:

DDQN performs significantly better than standard DQN under stochastic engine failures. This behavior is **fully consistent with the theoretical advantage of DDQN over DQN**.

- **DQN - Modified (Red):** Reaches a peak average reward of $\approx +130$ around episode 450, but suffers a severe **policy collapse** in late training, plummeting back down to $\approx +40$ by episode 600.
- **DDQN - Modified (Purple):** Recovers strongly from an early performance dip (≈ -150) and climbs steadily to finish near $+180$ by episode 600, maintaining an upward trend without collapsing.

Theoretical Explanation:

Standard DQN uses a single greedy operator to select and evaluate actions:

$$y_t^{\text{DQN}} = r_t + \gamma \max_{a'} Q_{\text{target}}(s_{t+1}, a')$$

Under environment stochasticity (misfires), noisy Q-value estimates lead to severe **maximization bias** ($E[\max(X)] \geq \max(E[X])$). Standard DQN overestimates sub-optimal or risky thruster actions, rendering the policy unstable and prone to collapse.

DDQN eliminates this bias by decoupling action selection from action evaluation:

$$y_t^{\text{DDQN}} = r_t + \gamma Q_{\text{target}}\left(s_{t+1}, \arg\max_{a'} Q_{\text{online}}(s_{t+1}, a')\right)$$

This structural separation prevents noise propagation, allowing DDQN to remain stable and converge in highly stochastic environments.

Question 5

Identify one limitation of your experimental setup and suggest one possible improvement.

Answer:

- **Limitation:** Across all 4 experimental arms, the strict safe landing success rate remained at **0%** throughout the entire 600 episodes (Plot 3). Evaluating over a single random seed ($N=1$) for only 600 episodes was insufficient for the agents to learn the fine-grained control needed to satisfy all 5 simultaneous landing criteria under an ϵ -greedy decay schedule.
- **Suggested Improvement:**
 - a. **Extend Training & Multi-Seed Runs:** Increase training budget from 600 to **1,500+ episodes** and average results across **5 to 10 independent random seeds** with shaded confidence intervals (mean $\pm$ standard error) to yield statistically rigorous conclusions.
 - b. **Curriculum Reward Shaping / Adaptive Exploration:** Implement an adaptive ϵ -decay schedule tied to touchdown velocity, or add a soft velocity-shaping term to incrementally guide the agent toward meeting the strict kinematic thresholds ($|v_x|, |v_y|, |\theta| < 0.10$).